

tem

Jason Guan

May 2017

1 Introduction

```
// Hello.java
import javax.swing.JApplet;
import java.awt.Graphics;

public class Hello extends JApplet {
    public void paintComponent(Graphics g) {
        g.drawString("Hello, world!", 65, 95);
    }
}
```




Figure 1: The Universe

2 DNN - Udacity Section 2 lecture 9

2.1 Multilayer Neural Network

ReLU

2.2 Practice TensorFlow ReLUs

```
// python
hidden_layer = tf.add(tf.matmul(features, hidden_weights), hidden_biases)
hidden_layer = tf.nn.relu(hidden_layer)
output = tf.add(tf.matmul(hidden_layer, output_weights), output_biases)
```

2.3 Tensorflow DNN

```
# input the dataset
from tensorflow.examples.tutorials.mnist import input_data
mnist = input_data.read_data_sets(".", one_hot=True, reshape=False)

# hyperparameter
import tensorflow as tf

learning_rate = 0.001
training_epochs = 20 # time of iteration
batch_size = 128 # lower this size if your memory is not enough
display_step = 1

n_input = 784 # MNIST data input, image size is 28*28
n_classes = 10 # MNIST total classes, total classes is 0-9 digits

# parameter in hidden_layer, which is the width of the layer
n_hidden_layer = 256

# weights and biases

weights = {
    'hidden_layer':tf.Variable(tf.random_normal((n_input,
        n_hidden_layer))),
    'out':tf.Variable(tf.random_normal((n_hidden_layer, n_classes)))
}
biases = {
    'hidden_layer':tf.Variable(tf.zeros(n_hidden_layer)),
    # tf.zeros should always take a vector or matrix as input? can a
    # scalar work
    'out':tf.Variable(tf.zeros(n_classes))
}

# input part
x = tf.placeholder('float', [None, 28, 28, 1]) # why there is always a
None
y = tf.placeholder('float', [None, n_classes])
```

```

x_flat = tf.reshape(x, [-1, n_input]) # convert a 28 * 28 matrix to a
    784 * 1 vector

# multilayer inception
layer_1 = tf.add(tf.matmul(x_flat, weights['hidden_layer']),
    biases['hidden_layer'])
layer_1 = tf.add(tf.matmul(layer_1, weights['out']), biases['out'])

logits = tf.add(tf.matmul(layer_1, weights['out']), biases['out'])

# Optimizer
cost =
    tf.reduce_mean(tf.nn.softmax_cross_entropy_with_logits(logits=logits,
        labels=y))
optimizer =
    tf.train.GradientDescentOptimizer(learning_rate=learning_rate).minimize(cost)

# Session
init = tf.global_variables_initializer()

with tf.Session() as sess:
    sess.run(init)

    for epoch in range(training_epochs):
        total_batch = int(mnist.train.num_examples/batch_size)
        for i in range(total_batch):
            # the method is provided by mnist
            batch_x, batch_y = mnist.train.next_batch(batch_size)
            # but how do they feed the dict x and y to optimizer?
            sess.run(optimizer, feed_dict={x: batch_x, y: batch_y})

```

2.4 About the Hidden Layer: why deeper not wider

1. Deep structure offers less parameters to estimate.
2. Deep structure is compatible with the abstraction learning pattern of image recognition.

2.5 Save Tensorflow Model and Restore it

```

# remove previous tensors and operations
tf.reset_default_graph()

from tensorflow.examples.tutorials.mnist import input_data
import numpy as np

# Hyperparameter
learning_rate = 0.001

```

```

n_input = 784
n_classes = 10

# Load dataset
mnist = input_data.read_data_sets('.', one_hot=True)

# Features and labels
features = tf.placeholder(tf.float32, [None, n_input])
labels = tf.placeholder(tf.float32, [None, n_classes])

# Weights and biases
# it is possible to trigger InvalidArgumentError: Assign requires shapes
# of both tensors to match.
# the name of tensor is better to be set explicitly
weights = tf.Variable(tf.random_normal([n_input, n_classes]),
                      name='weights')
biases = tf.Variable(tf.zeros([n_classes]), name='biases')

# Loss and optimizer
cost =
    tf.reduce_mean(tf.nn.softmax_cross_entropy_with_logits(logits=logits,
                                                            labels=labels))
optimizer =
    tf.train.GradientDescentOptimizer(learning_rate=learning_rate).minimize(cost)

# Calculate accuracy
correct_prediction = tf.equal(tf.argmax(logits, 1), tf.argmax(labels, 1))
accuracy = tf.reduce_mean(tf.cast(correct_prediction, tf.float32))

# Train model and save weights
import math

save_file = './train_model.ckpt'
batch_size = 128
n_epochs = 100

saver = tf.train.Saver()

with tf.Session() as sess:
    sess.run(tf.global_variables_initializer())

    # Loop over all batches
    for epoch in range(n_epochs):
        total_batch = math.ceil(mnist.train.num_examples / batch_size)

        for i in range(total_batch):
            batch_features, batch_labels =
                mnist.train.next_batch(batch_size)
            sess.run(optimizer, feed_dict={features: batch_features,
                                           labels: batch_labels})

```

```

# Print status for every 10 epochs
if epoch % 10 == 0:
    valid_accuracy = sess.run(
        accuracy,
        feed_dict={features: mnist.validation.images, labels:
                    mnist.validation.labels})
    print('Epoch {:<3} - Validation Accuracy: {}'.format(epoch,
        valid_accuracy))

saver.save(sess, save_file)
print('Trained Model Saved.')

saver = tf.train.Saver()

with tf.Session() as sess:
    saver.restore(sess, save_file)
    test_accuracy = sess.run(accuracy, feed_dict={features:
        mnist.test.images, labels:mnist.test.labels})

print('Test Accuracy:{}'.format(test_accuracy))

```

2.6 Regularization

1. Early Termination Stop to train as soon as the validation set performance begins to slide down.
2. L2 Regularization

$$L' = L + \beta \frac{1}{2} ||\omega||_2^2$$

3 Conclusion

“I always thought something was fundamentally wrong with the universe” [?]